

Impact Analysis of Botnet Infection on Networked Systems using Timed Automata

SYSC 5500
Group 8

Alvi Jawad
Luke Newton

Outline

Preliminaries

- Botnets in the IoT Network
- Timed Automata & UPPAAL

Modeling

- Overview of the Mirai Botnet
- Modeling The Botnet Infrastructure
- Modeling devices in the network

Simulations

- The Concrete Simulator

Results

- Data Collection
- Device Infection Rate
- Network Traffic Generation
- Next steps!

What are Botnets?

Botnets

- A network of **compromised devices**
- Compromised devices are called **bots**
- Bots infect other **vulnerable hosts** in the network
- Infection rarely requires **user interaction**



The Internet of Things

- **Vulnerable** networked infrastructure
- IoT devices
 - have **low** computational resources
 - are **poorly** secured
 - have **poor** maintenance
- IoT devices can generate **huge amounts of attack traffic** in a botnet

Why use Timed Automata?

Timed Automata

- Finite State Machines **extended with Clock variables**
- A **timed automaton** is a **6-tuple** (L, l_0, C, A, E, I)
 - L is a set of locations
 - $l_0 \in L$ is the initial location
 - C is the set of clocks
 - A is a set of actions, co-actions, and the internal τ -action
 - $E \subseteq L \times A \times B(C) \times 2^C \times L$ is a set of edges
 - $I : L \rightarrow B(C)$ assigns invariants to locations

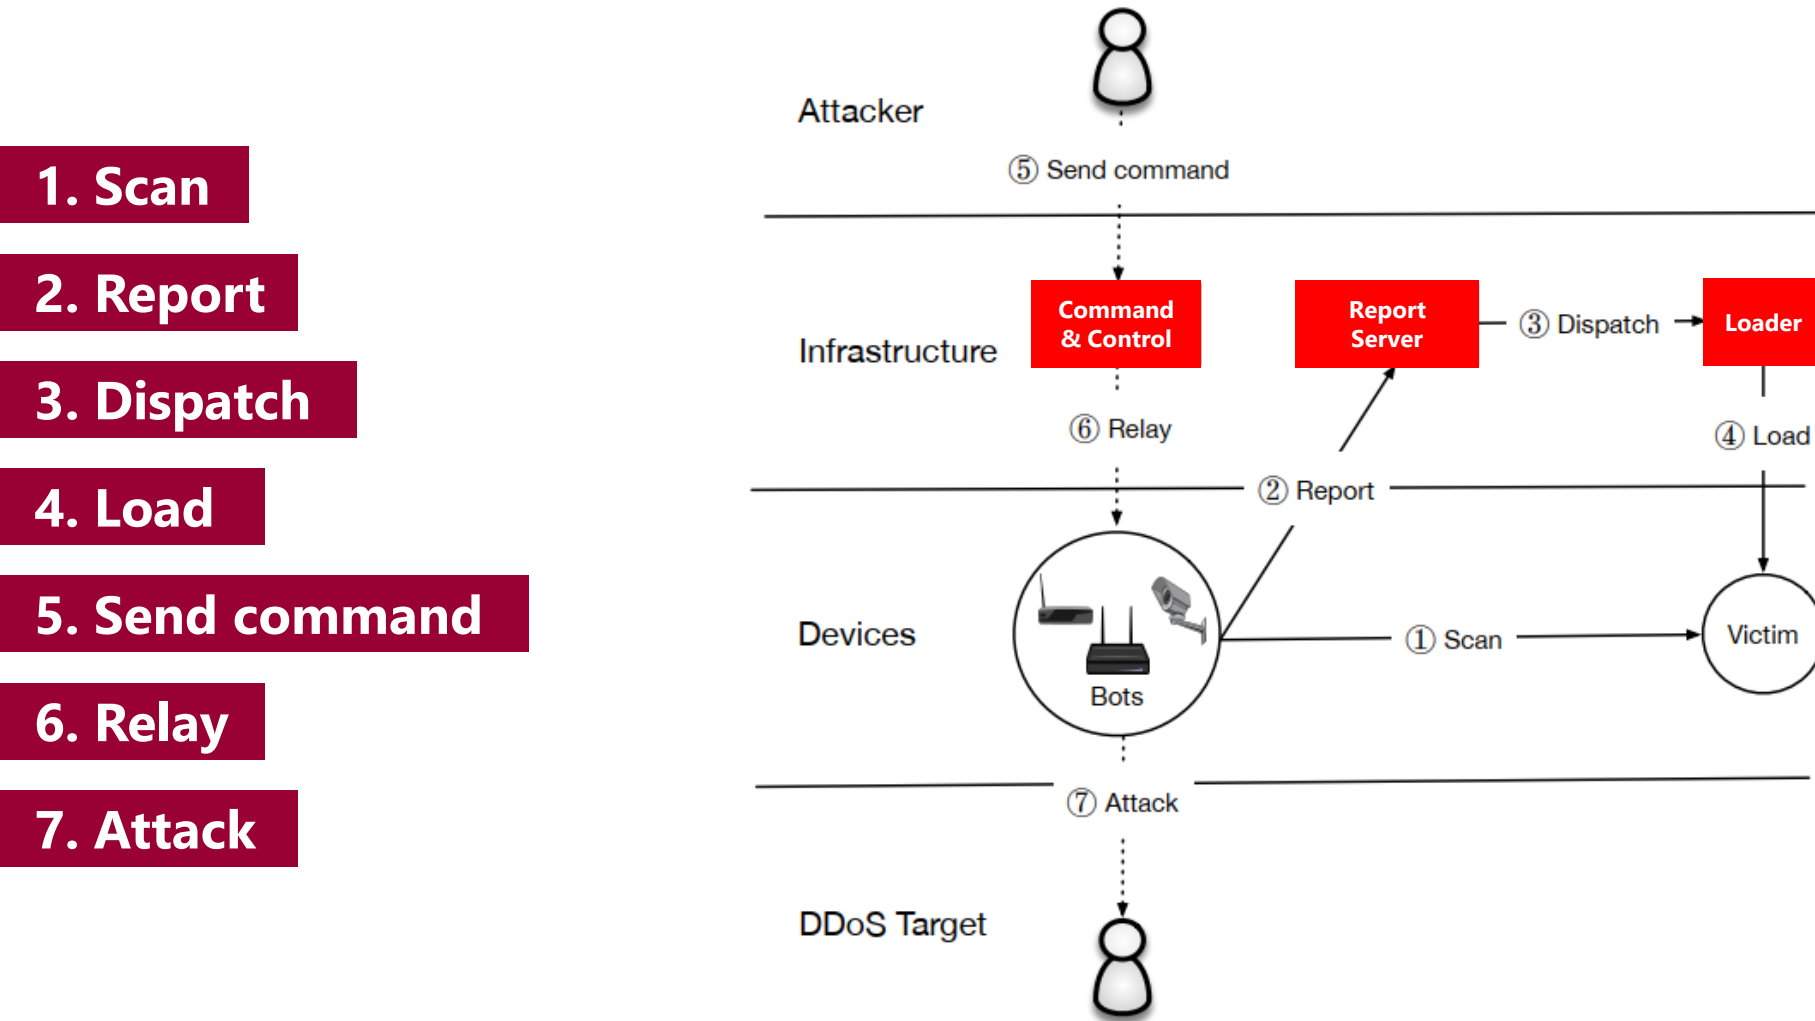
PROS

- Timed representation
- Hybrid-view of the system
- State Detection
- Behavior Prediction
- Formal verification

UPPAAL

- Verification of **real-time systems** modeled as networks of Timed Automata
- Query language: **TCTL** (Timed Computation Tree Logic)

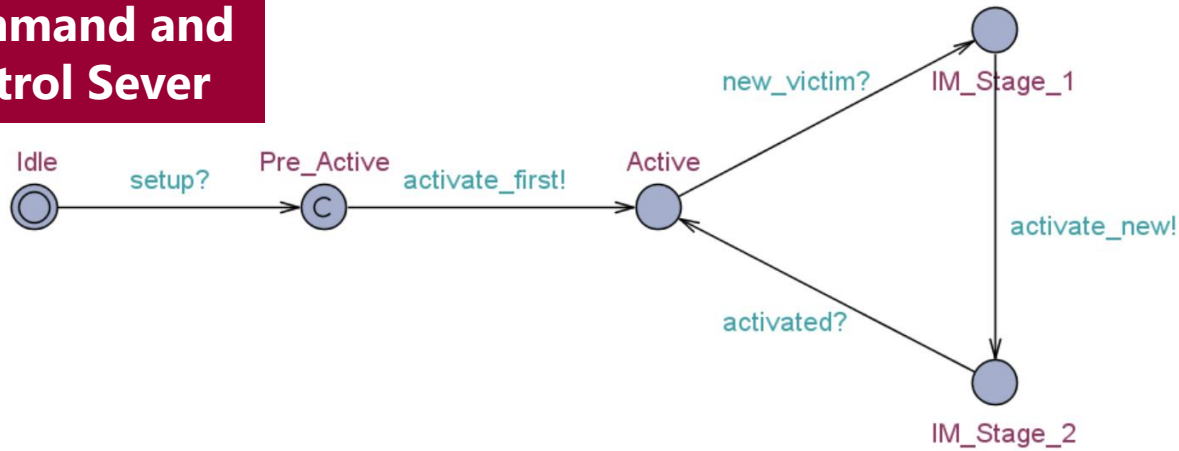
Overview of the Mirai Botnet



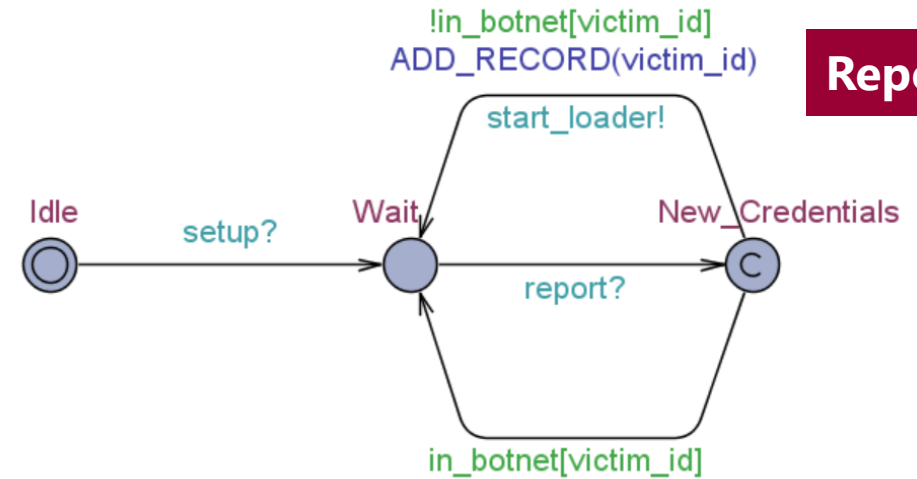
M. Antonakakis, T. April, M. Bailey, M. Bernhard, E. Bursztein, J. Cochran, Z. Durumeric, J. A. Halderman, L. Invernizzi, M. Kallitsis, D. Kumar, C. Lever, Z. Ma, J. Mason, D. Menscher, C. Seaman, N. Sullivan, K. Thomas, and Y. Zhou, "Understanding the mirai botnet," in *26th USENIX Security Symposium (USENIX Security 17)*, (Vancouver, BC), pp. 1093–1110, USENIX Association, Aug. 2017.

Modeling the Botnet Infrastructure


Command and Control Sever



Report Server

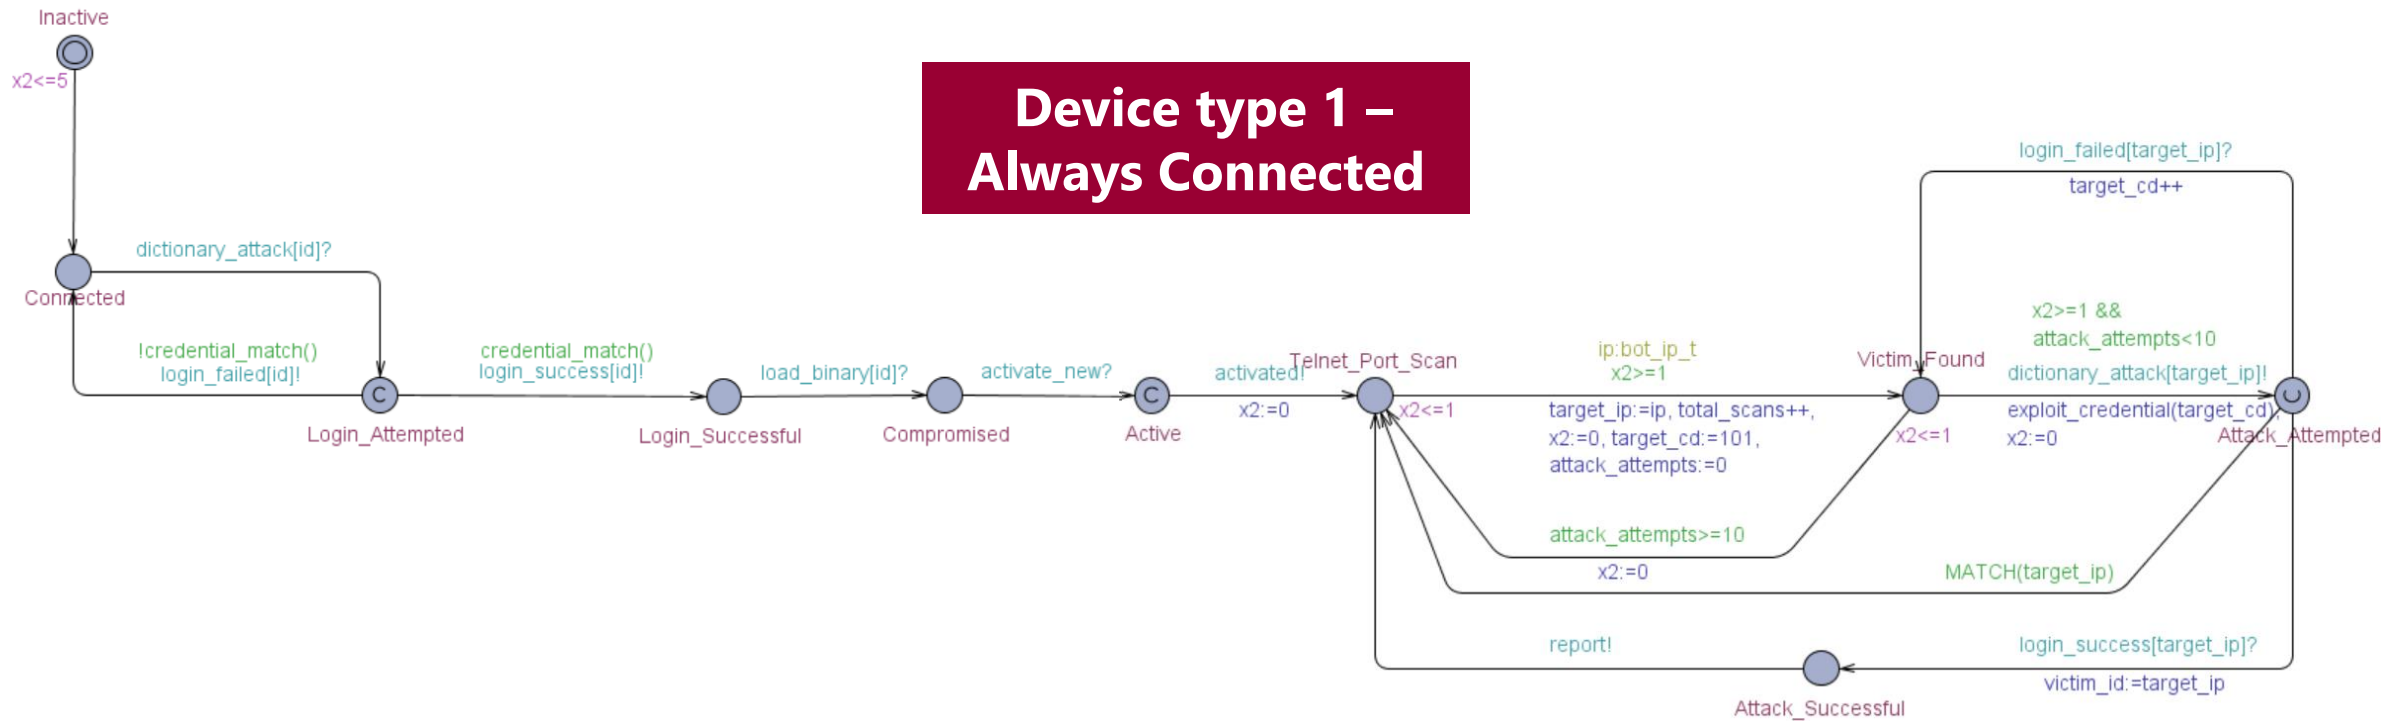


Loader



Modeling the Devices

Activation



Scan and Report

The Concrete Simulator – Initial State

State Variables

```
<Global variables>
total_scans = 0
total_attempts = 0
current_number_bots = 0
victim_id = 1
in_botnet = {0,0,0,0,0,0}
[0] = 0
[1] = 0
[2] = 0
[3] = 0
[4] = 0
[5] = 0
credential_attempt = {101,101,101,101,101,101}
[0] = 101
[1] = 101
[2] = 101
[3] = 101
[4] = 101
[5] = 101
t(0) = 0
total_time = 0.000000
```

```
BOT_DF
attack_attempts = 0
target_ip = 1
target_cd = 101
x1 = 0.000000

BOT1
attack_attempts = 0
target_ip = 1
target_cd = 101
x2 = 0.000000

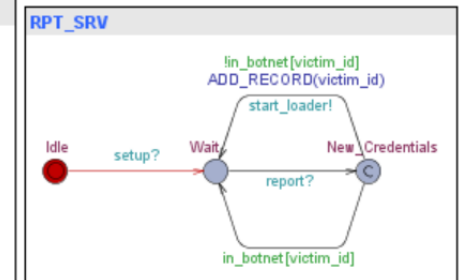
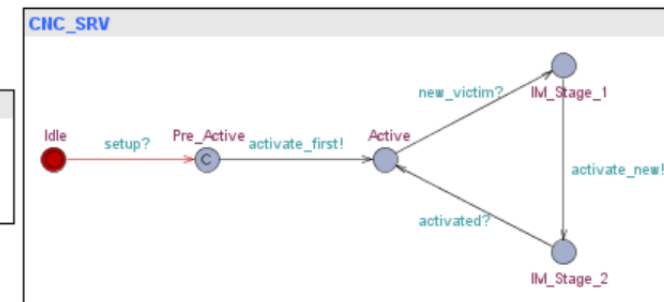


BOT2
attack_attempts = 0
target_ip = 1
target_cd = 101
x2 = 0.000000

BOT3
attack_attempts = 0
target_ip = 1
target_cd = 101
x2 = 0.000000

BOT4
attack_attempts = 0
target_ip = 1
target_cd = 101
x2 = 0.000000

BOT5
attack_attempts = 0
target_ip = 1
target_cd = 101
x2 = 0.000000
```

Choosing Transitions



The Concrete Simulator – in Action

State at 18 minutes

BOT_DF
attack_attempts = 1
target_ip = 2
target_cd = 102
x1 = 1.000000

BOT1
attack_attempts = 0
target_ip = 1
target_cd = 101
x2 = 18.000000

BOT2
attack_attempts = 0
target_ip = 1
target_cd = 101
x2 = 18.000000


BOT3
attack_attempts = 1
target_ip = 5
target_cd = 101
x2 = 0.000000

BOT4
attack_attempts = 6
target_ip = 2
target_cd = 107
x2 = 0.863843

BOT5
attack_attempts = 0
target_ip = 1
target_cd = 101
x2 = 18.000000

<Global variables>
total_scans = 7
total_attempts = 16
current_number_bots = 2
victim_id = 3
in_botnet = {1,0,0,1,1,0}
[0] = 1
[1] = 0
[2] = 0
[3] = 1
[4] = 1
[5] = 0
credential_attempt = {101,101,106,101,101,101}
[0] = 101
[1] = 101
[2] = 106
[3] = 101
[4] = 101
[5] = 101
t(0) = 0
total time = 18.000000

Gantt Chart



State at 29 minutes

BOT_DF
attack_attempts = 1
target_ip = 3
target_cd = 101
x1 = 0.000000

BOT1
attack_attempts = 0
target_ip = 4
target_cd = 101
x2 = 0.863843

BOT2
attack_attempts = 0
target_ip = 5
target_cd = 101
x2 = 0.000000

BOT3
attack_attempts = 1
target_ip = 1
target_cd = 101
x2 = 0.000000

BOT4
attack_attempts = 1
target_ip = 3
target_cd = 101
x2 = 0.863843

BOT5
attack_attempts = 0
target_ip = 1
target_cd = 101
x2 = 0.000000

<Global variables>
total_scans = 26
total_attempts = 46
current_number_bots = 5
victim_id = 5
in_botnet = {1,1,1,1,1,1}
[0] = 1
[1] = 1
[2] = 1
[3] = 1
[4] = 1
[5] = 1
credential_attempt = {101,101,101,101,101,101}
[0] = 101
[1] = 101
[2] = 101
[3] = 101
[4] = 101
[5] = 101
t(0) = 0
total time = 29.000000

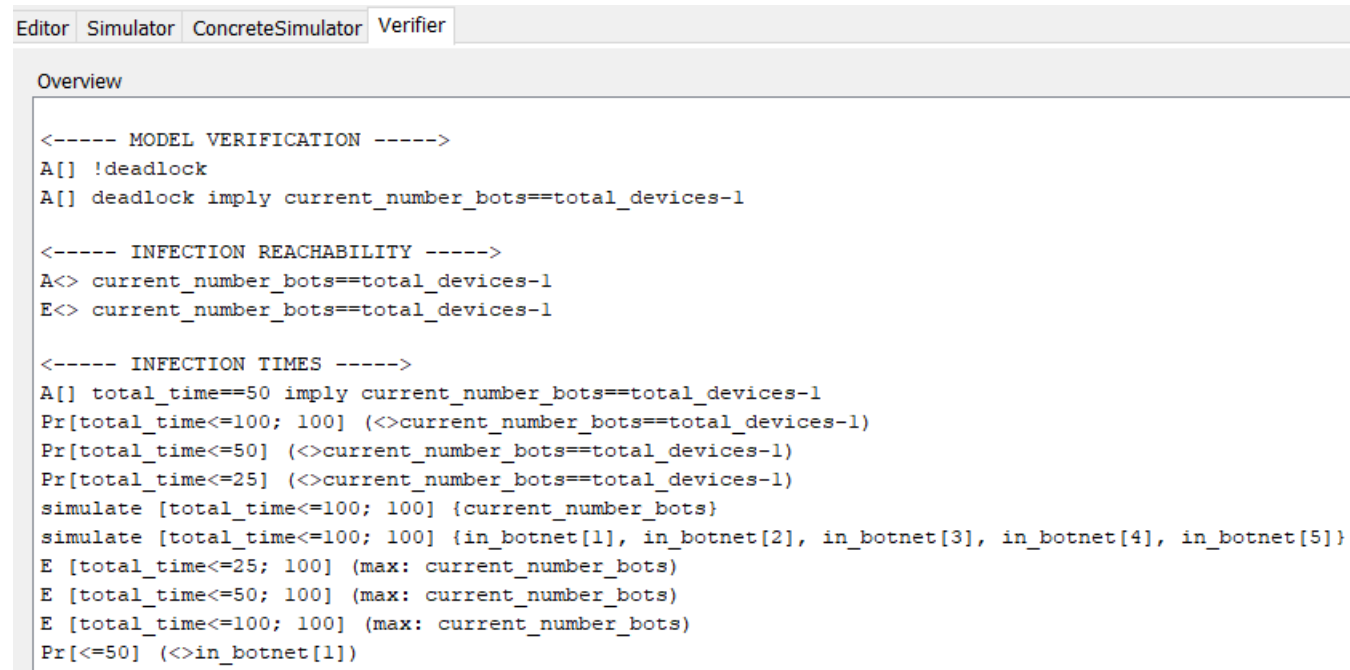
How do we collect data?

Concrete Simulator

- Run individual simulations and log data yourself
- See each run in **great detail**
- **Slow, manual** process

Verifier

- Formally check properties with **temporal logic**
- **Exhaustive** state space exploration
- Beware of **state space explosion!**
- Automate several simulation runs
- **Graph** frequencies and probability distributions
- Specify required **confidence levels**
- **Export data** for further analysis



The screenshot shows the Verifier application window with tabs for Editor, Simulator, ConcreteSimulator, and Verifier. The Verifier tab is active, displaying an 'Overview' section with model verification results. The results are organized into sections: MODEL VERIFICATION, INFECTION REACHABILITY, and INFECTION TIMES. The MODEL VERIFICATION section shows a property A[] !deadlock and its verification result A[] deadlock imply current_number_bots==total_devices-1. The INFECTION REACHABILITY section shows a property A<> current_number_bots==total_devices-1 and its verification result E<> current_number_bots==total_devices-1. The INFECTION TIMES section shows a property A[] total_time==50 imply current_number_bots==total_devices-1 and its verification result Pr[total_time<=50; 100] (<>current_number_bots==total_devices-1). The INFECTION TIMES section also shows a property Pr[total_time<=25; 100] (<>current_number_bots==total_devices-1) and its verification result Pr[total_time<=25; 100] (<>current_number_bots==total_devices-1). The INFECTION TIMES section also shows a property simulate [total_time<=100; 100] {current_number_bots} and its verification result E [total_time<=25; 100] (max: current_number_bots). The INFECTION TIMES section also shows a property simulate [total_time<=100; 100] {in_botnet[1], in_botnet[2], in_botnet[3], in_botnet[4], in_botnet[5]} and its verification result E [total_time<=50; 100] (max: current_number_bots). The INFECTION TIMES section also shows a property E [total_time<=100; 100] (max: current_number_bots) and its verification result E [total_time<=100; 100] (max: current_number_bots). The INFECTION TIMES section also shows a property Pr[<=50] (<>in_botnet[1]) and its verification result Pr[<=50] (<>in_botnet[1]).

```
<----- MODEL VERIFICATION ----->
A[] !deadlock
A[] deadlock imply current_number_bots==total_devices-1

<----- INFECTION REACHABILITY ----->
A<> current_number_bots==total_devices-1
E<> current_number_bots==total_devices-1


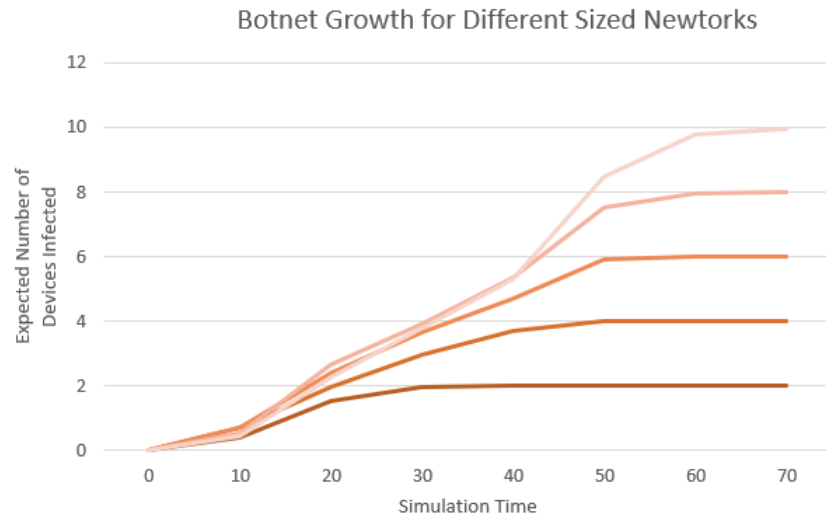
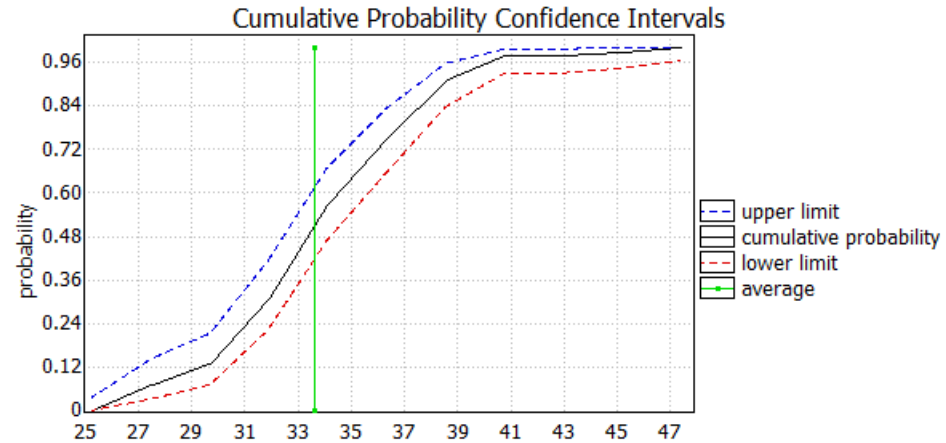
<----- INFECTION TIMES ----->
A[] total_time==50 imply current_number_bots==total_devices-1
Pr[total_time<=100; 100] (<>current_number_bots==total_devices-1)
Pr[total_time<=50; 100] (<>current_number_bots==total_devices-1)
Pr[total_time<=25; 100] (<>current_number_bots==total_devices-1)
simulate [total_time<=100; 100] {current_number_bots}
simulate [total_time<=100; 100] {in_botnet[1], in_botnet[2], in_botnet[3], in_botnet[4], in_botnet[5]}
E [total_time<=25; 100] (max: current_number_bots)
E [total_time<=50; 100] (max: current_number_bots)
E [total_time<=100; 100] (max: current_number_bots)
Pr[<=50] (<>in_botnet[1])
```

For now, we mainly produce results for a network of 5 devices

How quickly are devices infected?

Expected devices infected

- After 10 time units: 0.59 ± 0.0981
- After 25 time units: 2.93 ± 0.185
- After 50 time units: 4.99 ± 0.0198
- After 100 time units: 5 ± 0



Summary

- The model accurately depicts **expected botnet growth**
- Network size does not have a huge impact on expected time to infect all devices (so far)

Is network propagation “random”?

Devices are randomly targeted

- The **order** devices are infected **should not matter**
- Infection time distributions should be similar for each device

Kolmogorov-Smirnov Test

- Determine if two empirical distributions are sampled from the same distribution
- With confidence level 0.05 and 100 samples, **threshold value for the test is 0.192**
- Pairwise comparisons

What does this mean?

- **Location on the botnet password list significantly affects infection time**

Devices with different passwords

Device	2	3	4	5
1	0.37	0.23	0.23	0.29
2		0.29	0.32	0.17
3			0.10	0.18
4				0.28


All devices with same password:

Device	2	3	4	5
1	0.06	0.11	0.09	0.12
2		0.11	0.09	0.11
3			0.07	0.08
4				0.11

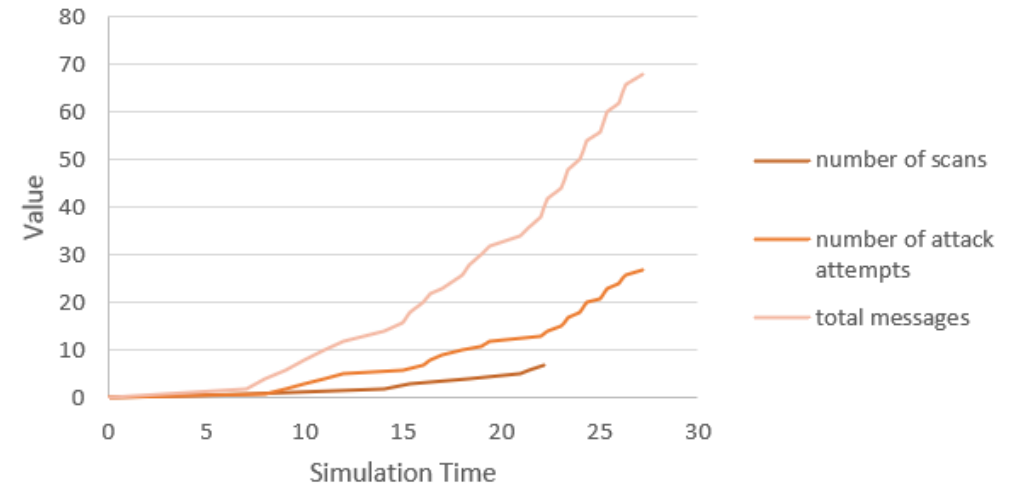
How much network traffic is generated?

- Messages sent by a botnet use resources that would otherwise be used for legitimate network traffic
- Expected number of **port scans**: 16 ± 12.5
- Expected number of **password attempts**: 38.5 ± 15.4
- Expected **total messages**: 109 ± 54.4

Expected Number of Messages Required to Infect Entire Network



Network Traffic Generated by Botnet Propagation



Summary

- Total **traffic generated** by the botnet **increases exponentially** over time
- Required **messages to infect whole network** **increases linearly** with network size

What comes next?

Rebooting and Patching

- Rebooting a device **clears any infection**
- Consider device types that reboot **periodically** and **following some distribution**
- Model **reinfectivity** of devices after patching and subsequent reboots

Network Stability

- Adjust timing on transitions

Password Scheme

- Widen the range of possible credentials a device can have

New Questions

- With rebooting and patching, is it possible to infect all devices?
- With rebooting, does the botnet size reach a steady state?
- How many times will a device be infected before the whole network is infected?
- How do these new considerations affect the amount of network traffic produced?

Thank You!

Questions?